FYP 2 Progress Report — Week 2 start (2026-10-12)

Project: Apa News Hari Tok? (Telegram news bot for Sarawak local news) **Author:** Vincent Lee Cheng Hong (BCS24020018) **Period covered:** 2026-10-11 → 2026-10-12 **Status:** SV has **approved the milestone-aligned FYP2 Gantt chart** (delivered in `Oct_Revised_FYP1_2 Gantt Chart.pdf`). First ~18 h of durable shadow evidence collected on Fly.io. All implementation bars through W7 of the approved chart are done or nearly done; software completion targeted by end of Week 2 (2026-10-18). **Code:** `main` @ `be8ac2f` deployed as Fly release v113

---

1. Gantt chart approved (2026-10-12)

Supervisor approved the revised FYP2 chart, now aligned to the official Milestones Verification Form (W1 plan review, W3 dev in progress, W4 completed implementation, W6 deployment, W7 test planning, W8 data collected, W11 Ch4 discussion, W12 final submission). Chart remains the *conservative plan*; actual progress runs ahead of it (see §3).

2. Shadow-mode evidence — first ~18 hours on Fly.io (2026-10-11 13:43Z → 10-12)

Artifact: `utility_alert_shadow_log` table (Supabase) + `flyctl logs`.

- **Stability**: polls tick every 10 min (~17–21 s each), zero job errors, zero source failures, zero shadow-log write failures.
- **Telegram source watch item CLOSED**: `swbnews` is being fetched from Fly's datacenter IP (3 posts/poll in current window) — the session works from the server; no auth errors at any point.
- **2 new real alerts detected**: 10-11 13:20Z (117 would-be deliveries) and 10-12 07:03Z (new JBALB/news alert, 132 would-be deliveries).
- **Quiet hours held**: the only would-deliver computation in the window landed at 15:03 Kuching time; nothing computed during 00:00–06:00 Asia/Kuching.
- **Cross-instance idempotency**: laptop → Fly handover deduped cleanly against the shared database.

Stage 1b criteria status: zero pushes [OK], stability [OK] (in progress), no cross-poll duplicates [OK], quiet hours [OK] (no counter-example yet), precision / rate-cap — pending a larger stored-alert sample.

3. Position vs the approved chart (start of Week 2)

Approved bar (week)	Actual
Project Plan Review & Refinement (W1)	[OK] done (chart approved 10-12)
Database Infrastructure Migration (W2)	[OK] done 10-05
Summary Length Consistency (W3)	[OK] done 10-05
RAG & Hybrid Retrieval (W4)	[OK] done 10-06
Semantic Dedup & Tuning Set (W4–5)	■ core done 10-06; human-labeled tuning set open
Utility Alert Integration (W5)	■ Stage 1 done 10-06; Stage 2 pending owner sign-off
NLP Benchmarking (W6–7)	[OK] done 10-12 (doc 07: llama3.1 local vs gpt-oss:20b cloud)
Deployment + Repository URL (W6–7)	[OK] done 10-11 (Fly v113, repo public)

Everything scheduled through W7 is complete except the benchmark — roughly 3–5 weeks ahead of the approved chart.

4. Open items to close the software phase (target: 2026-10-18)

1. **NLP benchmark (mimo)** — **done 10-12 as doc 07**, with a scope correction: the cloud arm is **gpt-oss:20b** (the model the bot actually falls back to), not MiMo (fallback_2, invalid key). Four-dimension framework executed with automatic metrics + cross-model LLM judging. Production-relevant finding: the 384-token `OLLAMA_SUMMARY_NUM_PREDICT` cap truncates the reasoning model to empty output — raise it before relying on the cloud fallback.
2. **Dedup human-labeled tuning set** — **complete 10-12**: 40 owner-labeled pairs (`benchmark/dedup_tuning_set_20261012.xlsx`) + second-signal fix implemented and validated (doc 03

§3.3–3.4; labeled-set F1 0.541 → 0.692; production A/B clean).

3. **Utility-alert Stage 2 checklist** (04 §6): owner sign-off from shadow evidence, `/settings` toggle for `wants_urgent_alerts`, incident-level clustering, restoration threading, then `UTILITY_ALERT_SHADOW_MODE=false`.

4. **Hygiene**: Supabase password rotation (Fly holds `DATABASE_URL`), decision on RSS 403 mitigation (Fly IPs blocked by See Hua / Borneo Post / FMT / TVSarawak — prefetch either moves to the laptop or coverage is documented as a limitation).

5. Cosmetic: `logging.basicConfig(level=INFO)` so the service `NEW`-alert detail lines reach the log stream.

Next report: Week 3 entry expected to show benchmark results + tuning set + Stage 2 decision, marking the software phase complete and shifting focus to the survey instrument and report writing.

Previous: 05-fly-shadow-deployment